

Deep Generative Modelling

Presenter:

Kurra Bavanya Choudhry

18075034
BTech CSE Part-IV

Core goal of Generative modelling

- Let's say we have an observed data D , for eg: images of cats.
- D can be thought of as a sample from probability distribution p_{data} (true distribution).
- The goal of Generative modelling is to find an approximate probability distribution p_{θ} which is as close as possible to p_{data} for an observed data D .
- Now we can generate unseen data by sampling from p_{θ} .



← Surprisingly, this is not an image of a real child!

This image was produced by the style based generator proposed by T Karras in 2018.

source: [Karras et al. \(2018\)](#)

2 Tasks: Learning and Inference (1/2)

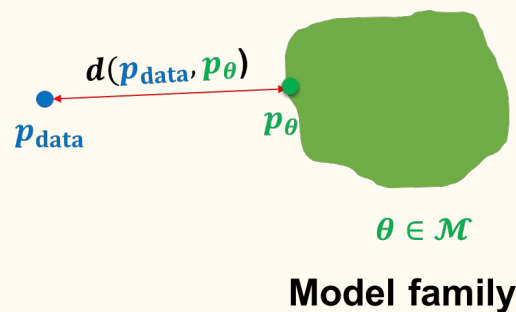
- **Learning (unsupervised and parametric):**

- p_θ is picked from a family of model distributions M .
- **Optimization problem:**
$$\min_{\theta \in M} d(p_{\text{data}}, p_\theta)$$

- where: $d(\cdot)$ is a notion of distance between probability distributions.

- **Three questions arise here:**

- How to represent model family M ?
- What is the objective function $d(\cdot)$?
- How do we minimize $d(\cdot)$?



2 Tasks: Learning and Inference (2/2)

- Inference

- Generative models are used for 3 fundamental inference tasks:

- Density estimation:

- Obtaining $p_{\theta}(\mathbf{x})$ for a given datapoint \mathbf{x} .

- Sampling:

- Generating unseen data $\mathbf{x}_{\text{new}} \sim p_{\theta}(\mathbf{x})$.

- Unsupervised representation extraction:

- Latent space feature representation vector $\mathbf{f}_{\text{vector}}$ for a given datapoint \mathbf{x} .

There are many deep generative models (DGMs) present in the market today and we choose a model suitable for our inference requirements.

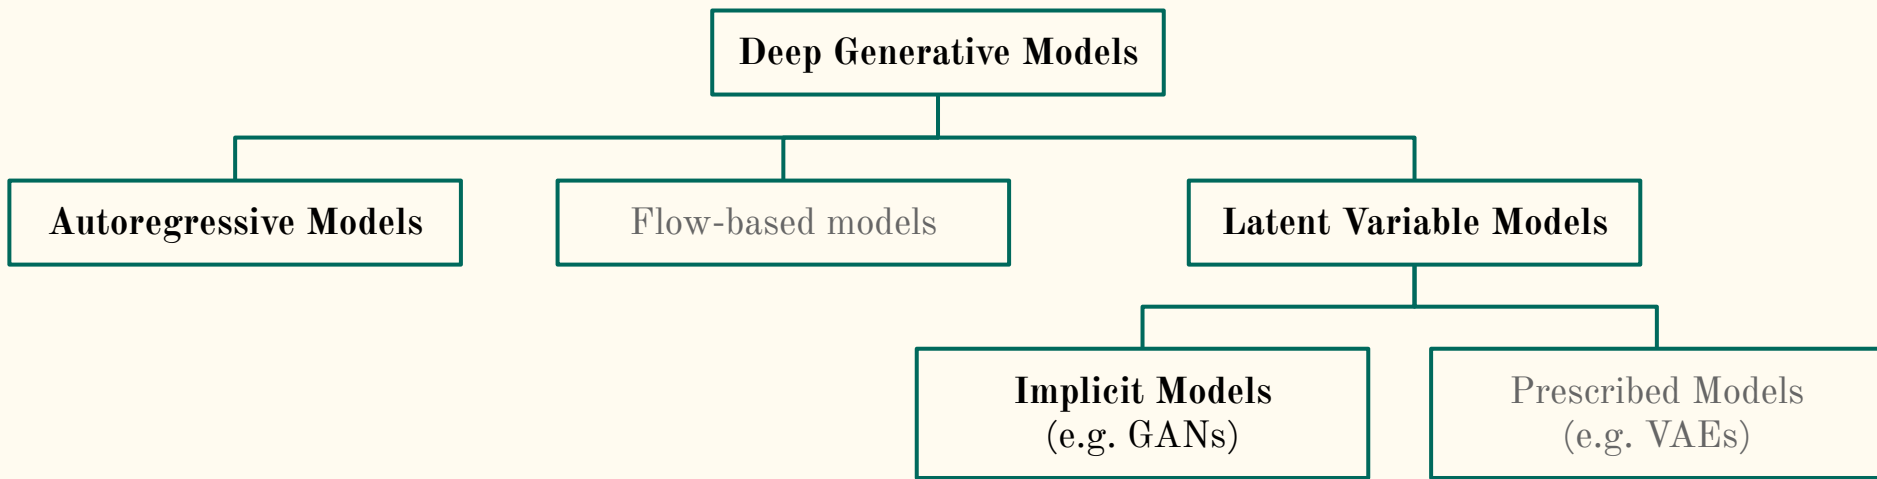
Generative approach to classification task

- **Problem:** Detect an image as cat or dog):
- **Solution:**
 - Obtain p_{θ} for D_{cat} and p_{θ_1} for D_{dog} .
 - where: $D_{cat} \cup D_{dog} = D$ & $D_{cat} \cap D_{dog} = \Phi$.
 - For test datapoint \mathbf{x}_{test} , find $p_{\theta}(\mathbf{x}_{test})$ and $p_{\theta_1}(\mathbf{x}_{test})$
(to find out which distribution is it most probably came from.)
- **How is it different from discriminative approach?**
 - Doesn't find a decision boundary that separates data points of different classes.
 - Estimates $P(Y|X)$ using Bayes Rule on $P(X|Y)$ i.e $p_{\theta}(\mathbf{x})$ and $P(Y)$, whereas discriminative models finds it directly from training data.

Family of Deep Generative Models (DGMs)

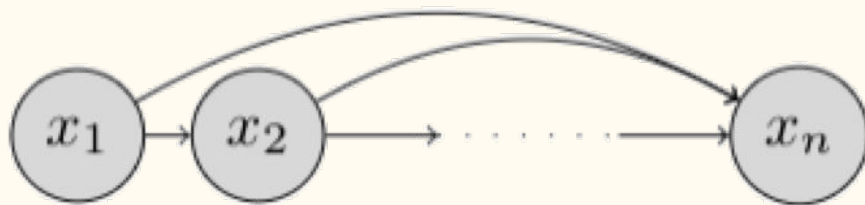
- "**Deep**" in DGM refer to multiple layers of neurons with hidden variables.
- Generative modelling can be done without using neural networks (NNs).
- But still NNs are widely used to parameterize generative models because they are flexible and powerful.

Deep Generative models can be expressed into the following three main groups:



Autoregressive Models (1/2)

- We fix an ordering of the variables x_1, x_2, \dots, x_n .
- The distribution for the i -th random variable depends on the values of all the preceding random variables in the chosen ordering x_1, x_2, \dots, x_{i-1} .
- Expressing the above assumption in graphical form:



(source: [website](#))

- $p_{\theta}(\mathbf{x}) = \prod_{i=1}^n p_{\theta}(\mathbf{x}_i | \mathbf{x}_{<i})$ where $\mathbf{x}_{<i} = [x_1, x_2, \dots, x_{i-1}]$

Autoregressive Models (2/2)

- KL divergence is taken as the objective function $d(\cdot)$.

$$\min_{\theta \in M} d_{\text{KL}}(p_{\text{data}}, p_{\theta}) = \max_{\theta \in M} (\sum_{\mathbf{x} \in D} \log(p_{\theta}(\mathbf{x})))$$

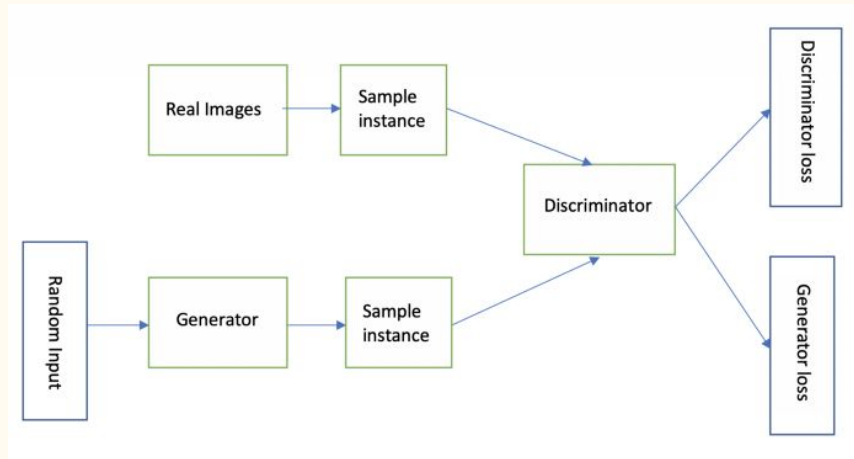
i.e we pick the model parameters $\theta \in M$ that maximize the log-probability of the observed data points in D .

- The model parameters $\theta \in M$ are optimized using mini-batch gradient descent.
- For density estimation and unseen data generation, sequential sampling procedure is followed.
- Hence both the inference tasks are expensive.
- These models do not learn unsupervised representation of data.
- Hence, they can't be used to get $\mathbf{f}_{\text{vector}}$ for a given datapoint.

Generative Adversarial Networks (GANs)

- It has two parts:
 - **Generator**
 - Learns to generate plausible data.
 - Its output become negative training examples for the discriminator.
 - **Discriminator**
 - Learns to distinguish the generator's fake data from real data.
 - Penalizes the generator for producing implausible results.

GAN Architecture



source: [medium post](#)

Training in GANs

- **at Discriminator:**

- Connects to two loss functions: Generator and Discriminator loss.
- Ignores the generator loss and just uses the discriminator loss.
- The discriminator loss penalizes the discriminator for misclassifying a real instance as fake or a fake instance as real.
- Network's weights are updated through backpropagation.

- **at Generator:**

- Takes random noise sampled from any probability distribution as input.
- This noise is transformed into meaningful output.
- The gradient of the Generator loss is used to change the weights of only the generator.

Real life applications

- **Data Generation**

- Art creation
- Code generation
- Text generation
- Audio synthesis
- Text to image generation
- etc.



New York City in a thunderstorm (cc12m_1_cfg)



source: [post link](#), generated using [cc12m_1_cfg model](#) developed by [@RiversHaveWings](#)

- **Measuring the plausibility of data**

- Anomaly detection
 - If $p_{\theta}(\mathbf{x}^1) \lll p_{\theta}(\mathbf{x})$ for $\mathbf{x} \in D$, then \mathbf{x}^1 is an outlier.
- Classification (*it was discussed in previous slides*)
- etc.

THANK
YOU

—